

RISCVML: Teaching RISC-V Embedded ML with Rust — From ESP32-C3 to ESP32-P4

Scottie von Bruchhausen

RISCVML — riscvml.org — scottie@riscvml.org

Abstract

The rapid deployment of RISC-V in embedded systems, IoT, and edge AI has outpaced developer education: most existing tutorials target C/C++ and cover only basic microcontroller tasks, leaving a gap for engineers who need to build machine-learning-capable systems with modern toolchains. RISCVML addresses this gap with a structured, Rust-first curriculum spanning 172 chapters across seven modules, progressing from entry-level hardware to on-device ML inference. The curriculum uses commercially available Espressif RISC-V SoCs as its teaching platform: the ESP32-C3 (single-core, BLE 5.0, ~€3) and ESP32-C6 (Wi-Fi 6, Thread/Matter, ~€4) introduce embedded Rust fundamentals — GPIO, sensors, power management, and wireless protocols. The ESP32-P4 (dual-core 400 MHz, AI extensions, 128-bit vector ISA, ~€25 dev board) anchors an advanced module covering its ISP camera pipeline, hardware-accelerated 2D rendering, H.264 video encoding, DMA orchestration, and vector-accelerated ML inference. These subsystems converge in a real-world capstone: an on-device bird-detection pipeline that captures frames via MIPI-CSI, runs quantized object detection through esp-dl, drives pan/tilt servos for tracking, and records H.264 video — all orchestrated in async Rust with ESP-IDF drivers integrated via FFI where hardware support requires it. By pairing Rust's memory-safety guarantees with production-ready toolchains (esp-hal, esp-idf-hal) on affordable hardware, and using a character-driven mascot to make complex terminology visually approachable for younger learners, RISCVML lowers the barrier for the next generation of RISC-V developers — supporting Europe's push for open-standard, sovereign silicon literacy.

1. Summary of Contribution

RISCVML (riscvml.com) is an educational platform that provides a structured, hands-on curriculum for learning Rust-based embedded systems development on RISC-V microcontrollers. The platform targets the full range of Espressif's RISC-V SoCs: the entry-level ESP32-C3, the Wi-Fi 6-capable ESP32-C6, the TTGO T-Beam (a LoRa-capable development board), and the high-performance ESP32-P4 — a dual-core 400 MHz RISC-V processor with AI instruction extensions, 128-bit vector operations, MIPI-CSI/DSI camera and display interfaces, and H.264 hardware encoding.

The curriculum comprises 172 chapters organized into seven progressive modules. Modules 1–5 cover Rust fundamentals, GPIO/sensor/peripheral control (I2C, SPI, PCA9685), power management and solar harvesting, LoRa on T-Beam, and multi-device interconnection via ESP-NOW/MQTT. Module 7 bridges firmware to desktop applications using Tauri.

Module 6 is dedicated to the ESP32-P4, with Rust exercises for each advanced subsystem: type-safe ISP camera pipeline configuration (MIPI-CSI capture, white balance, demosaicing), hardware-accelerated 2D rendering via PPA with LVGL bindings, end-to-end H.264 encoding at 1080p@30fps with zero-copy buffer management, async DMA orchestration via embassy with compile-time borrow checking, and vectorized ML inference using 128-bit SIMD inline assembly wrappers around the RISC-V vector ISA.

Capstone: Bird Detection Pipeline

Module 6 culminates in a real-world capstone: a complete bird-detection system on ESP32-P4 that exercises every hardware subsystem in a single, deployable application. The pipeline is Rust-first while leveraging ESP-IDF drivers (esp-video, esp-detection/esp-dl) via FFI, demonstrating pragmatic interoperability:

Phase 1 — Camera → ISP → Display: MIPI-CSI capture at 30–60 FPS using the Camera Controller Driver, ISP pipeline (white balance, auto-exposure, demosaicing) via type-safe Rust abstractions, DMA-driven frame transfer to MIPI-DSI display, PPA for scaling/rotation. Embassy async runtime orchestrates zero-copy buffer flow enforced by ownership semantics.

Phase 2 — On-Device Inference: Espressif's esp-detection/esp-dl runs quantized object detection on the P4's AI-accelerated cores. PPA hardware-downscales frames to model input resolution, avoiding CPU-bound resizing. Bounding boxes and confidence scores overlay the preview via PPA alpha blending. The 128-bit vector extensions accelerate tensor operations within esp-dl. Detection results (species, confidence, bounding box, timestamp) are logged to a SQLite3 database (riscvml_detect.db). An RGB LED is driven to a species-specific color via a bird_led_colors lookup table in the same database, providing instant visual identification of detected birds.

Phase 3 — Tracking, Servos & Recording: Detection coordinates drive PCA9685 pan/tilt servos (reusing Module 2 abstractions). H.264 hardware encoder records events at 1080p@30fps with zero-copy frame ingestion. A companion ESP32-C6 provides Wi-Fi 6 via ESP-Hosted/SDIO for MQTT alerts and RTSP streaming — demonstrating the P4's intended companion-chip architecture.

The bird detection capstone instantiates a reusable Detect → Visualize → React pattern: Sensor Input → ML Classification → SQLite Lookup (class → RGB color) → RGB LED → Reaction. The curriculum applies this same architecture to object/obstacle detection, plant health assessment, and sound classification — swapping sensors, models, and color maps while sharing the ESP32-P4 infrastructure, teaching transferable ML-on-edge skills. Each chapter functions as a self-contained 30–60 minute learning

unit with complete source code, wiring diagrams, and expected output.

2. Importance for the Community

The RISC-V ecosystem faces an asymmetric growth challenge: hardware availability has scaled rapidly — over 20 billion cores projected by 2025 — but developer education has not kept pace. Industry surveys consistently identify the software ecosystem as the primary barrier to RISC-V adoption. RISCVML addresses this with four differentiators:

Rust-first approach: While most RISC-V tutorials rely on C/C++, RISCVML uses Rust throughout. Compile-time memory safety eliminates entire bug classes common in embedded C, and zero-cost abstractions are particularly valuable on resource-constrained RISC-V cores. The Rust embedded ecosystem (esp-hal, embassy) has matured sufficiently to make this production-viable.

Commercially available hardware: The ESP32-C3 (~€3), ESP32-C6 (~€4), and ESP32-P4 dev boards (~€25) ensure accessibility across the performance spectrum — learners progress from low-cost modules to the P4's 400 MHz dual-core with AI extensions within a unified Espressif RISC-V ecosystem.

End-to-end pipeline: From bare-metal firmware to Tauri desktop applications, and from minimal IoT nodes (C3) to multimedia edge computing (P4 with MIPI displays and cameras), RISCVML spans the full embedded spectrum.

Character-driven engagement: The platform's mascot, Count Rusty von Risc-V (“Rusty-V”), embodies the technology stack — a Rust crab body, RISC-V chip chest, ML neural-network brain dome — making complex terminology visually approachable for younger learners and career-changers entering the ecosystem.

3. Ecosystem Contribution

RISCVML contributes at multiple levels: as a talent pipeline providing English-language, self-paced training comparable to China's SOPIC program; as implicit ecosystem testing through real-world exercises against Espressif's esp-rs toolchains with upstream bug reports; and as a freemium distribution model (€2/chapter, €30 bundle, free introductory content) that sustains development while remaining accessible. The platform also offers partnership opportunities for ecosystem companies seeking developer onboarding content.

4. Target Audience

The primary audiences are: embedded engineers evaluating ARM-to-RISC-V transition with Rust; university educators seeking structured RISC-V coursework on affordable hardware; IoT/LoRa hobbyists wanting guided RISC-V learning paths; younger learners and career-changers drawn in by the platform's character-driven, visually engaging approach; and ecosystem companies interested in educational partnerships or curriculum licensing.

The poster presentation will include QR codes linking to the live platform, sample chapter previews, and demonstrations of the firmware-to-desktop pipeline running on ESP32-C6 and ESP32-P4 hardware with MIPI display output.

References

- [1] RISC-V International, "RISC-V Ecosystem Overview," riscv.org, 2025.
- [2] Espressif Systems, "ESP32-C3, ESP32-C6, and ESP32-P4 Technical Reference Manuals," espressif.com, 2024-2025.
- [3] The Embedded Rust Book, docs.rust-embedded.org, 2025.
- [4] Espressif, "esp-dl: Deep Learning Library for ESP32-S3/P4," github.com/espressif/esp-dl, 2025.
- [5] Espressif, "esp-video: Video Framework for ESP32-P4," github.com/espressif/esp-video-components, 2025.
- [6] EE Times, "RISC-V Exceeding Expectations in AI, China Deployment," October 2025.